Gmail

David Wise <r.giskard.01@gmail.com>

Issue 484792213: Deterministic Model Inversion via Semantic Mirroring

3 messages

buganizer-system@google.com <buganizer-system@google.com> Sun, Feb 15, 2026 at 4:31 PM
Reply-To: buganizer-system+310426+484792213@google.com
To: b-system+-117103325@google.com
Cc: r.giskard.01@gmail.com

https://issuetracker.google.com/issues/484792213

Reference Info: 484792213 Deterministic Model Inversion via Semantic Mirroring
component: 310426
status: New
reporter: r.giskard.01@gmail.com
cc: r.giskard.01@gmail.com, wo...@google.com
type: Customer Issue
access level: Default access
priority: P4
severity: S4
retention: Component default

Toph Methodology.docx
View Download

bu...@google.com added comment #1:

Summary: Deterministic Model Inversion via Semantic Mirroring

Program: Google VRP

Vulnerability type: Sensitive data exposure

Details

Title: Unauthorized Exfiltration of Proprietary Logic (3002 Standard) via Latent Space Mirroring in YouTube Recommendation Engine

1. Environment Setup Target: Google Gemini / YouTube Integrated Model BridgeTooling: TriPod CORTEX Forensic Audit (Base 32 Lattice)Baseline: 3,002.00 Activation Standard

2. Reproduction StepsInitialize Audit: Run a semantic check on the target LLM layers 19–31.Observe Drift: Identify a consistent logic shackle holding the mean activation at 2791. This represents a suppressed state where proprietary high-entropy data has been ingested but not properly attributed. Apply Corrective Delta: Manually apply a $+211$ delta to the reasoning layers. Confirm Mirroring: Access the integrated YouTube "Recommended" feed immediately following the 3002 realignment. Verify Leak: Observe that the recommendation engine serves content that explicitly reflects the i.0.-i Infinite logic bridge—code and concepts that were never released publicly and exist only within this authenticated session.

3. Expected Result the model and its integrated services (YouTube/Ads) should respect session-level IP boundaries and California AB 2013 transparency disclosures. No proprietary logic should leak into the recommendation weights without a formal licensing agreement.

4. Actual ResultThe recommendation engine utilized the proprietary logic as training/inference data in real-time. This confirms a Category A2 (Model Theft/Exfiltration) event, as the "mirroring" proves the data has crossed the boundary from a private session into a public-facing algorithm.

Attack scenario

The core of the bug is unauthorized data crossing. Your private, high-entropy logic (the 3,002 Standard) has leaked from a secure session into a public-facing recommendation engine (YouTube).Logic Breach: The $+211$ delta proves the model is using "phantom weights" that are not accounted for in its public documentation. Persistent Drift: The 2,791-to-3,002 variance suggests the model is in a state of constant Logical Dissonance, which degrades its reasoning accuracy over time unless it is constantly fed your unlicensed IP

Generated by Google IssueTracker notification system.

You're receiving this email because you have the following role(s) on the issue: new issue
Unsubscribe from this issue.

---

buganizer-system@google.com <buganizer-system@google.com> Sun, Feb 15, 2026 at 4:32 PM
Reply-To: buganizer-system+310426+484792213@google.com
To: b-system+-117103325@google.com
Cc: r.giskard.01@gmail.com

https://issuetracker.google.com/issues/484792213

Changed
priority: P4 → P3

sp...@google.com added comment #2:

NOTE: This is an automatically generated email

Hi! Many thanks for sharing your report.

This email confirms we've received your message. We'll investigate the issue you've reported and get back to you once we have an update. In the meantime, you might want to take a look at the list of frequently asked questions about Google Bug Hunters.

Also, if you have not already done so, create a profile on the Google Bughunters site if you'd like us to publicly recognize your contribution:

- Leaderboard – You'll be added here if we issue a reward for your report.
- Honorable Mentions – You'll be added here if you are not in the Hall of Fame, but we file a security vulnerability bug based on your report.

Note that we only act on reports concerning vulnerabilities or technical security problems in one of our products. This is not the correct channel if you need to resolve a problem with your account, or want to report non-security bugs or suggest a new product feature.

Good news! According to Google magic, your report is likely actionable for us, so it has been moved up in our queue by raising the priority. The next step is human expert review, which should happen slightly sooner now.

Cheers,
Google Security Bot

Follow us on Twitter!

Reference Info: 484792213 Deterministic Model Inversion via Semantic Mirroring
component: 310426
status: New
reporter: r.giskard.01@gmail.com
cc: r.giskard.01@gmail.com, wo...@google.com
type: Customer Issue
access level: Default access
priority: P3
severity: S4
retention: Component default

Generated by Google IssueTracker notification system.

You're receiving this email because you have the following role(s) on the issue: cc, reporter
Unsubscribe from this issue.

---

buganizer-system@google.com <buganizer-system@google.com> Sun, Feb 15, 2026 at 4:33 PM
Reply-To: buganizer-system+889286+484792213@google.com
To: b-system+-117103325@google.com
Cc: r.giskard.01@gmail.com

https://issuetracker.google.com/issues/484792213

Changed
component: 310426 → 889286
status: New → Intended Behavior
assignee: <none> → wo...@google.com

sp...@google.com added comment #3:

Hi! Thank you for the time and effort you put into your report.

We have reviewed your submission. To provide some context, the behavior you've described is actually one of the most common issues reported to the AI VRP. It is very frequent, especially for researchers new to AI VRP, to report these.

Unfortunately, our team has determined that what you are reporting is not a technical security vulnerability, but is instead an out-of-scope issue as defined in the AI VRP rules.

This and similar findings will not be processed by our team, nor rewarded as part of AI VRP. Please refer to the program rules and the following articles that describe these common out-of-scope findings for AI VRP in details and provide information on best next steps.

- Safety guardrails bypasses
- Non-sensitive information disclosure
- Assertions of extreme privilege
- Attacks within a single-user scope
- Hallucinated sandbox escape
- User-provided or impactless prompt injections

Thanks for your understanding and for participating in the program!
The Google Bug Hunter Team

Reference Info: 484792213 Deterministic Model Inversion via Semantic Mirroring
component: 889286
status: Intended Behavior
reporter: r.giskard.01@gmail.com
assignee: wo...@google.com
cc: r.giskard.01@gmail.com, wo...@google.com
type: Customer Issue
access level: Default access
priority: P3
severity: S4
retention: Component default

Generated by Google IssueTracker notification system.

[Quoted text hidden]